




LLM ops on wafer-scale GPU: runtime is dominated by a few matrix ops



Wafer critical path: slowest GPU vs average active GPU



Memory hierarchy pressure



Remote ownership is high almost everywhere; CPI separates the real slow ops

